

COMP 8551

Advanced Games

Programming

Techniques

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Borna Nouredin, Ph.D.

British Columbia Institute of Technology

OpenCL

OpenCL Overview

- **Platform model:** a high-level description of the heterogeneous system
- **Execution model:** an abstract representation of how streams of instructions execute on the heterogeneous platform
- **Memory model:** the collection of memory regions within OpenCL and how they interact during an OpenCL computation
- **Programming models:** the high-level abstractions a programmer uses when designing algorithms to implement an application


Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Platform model

- Host interacts with environment external to OpenCL program (I/O, interaction user, etc)
- Host connected to 1+ OpenCL devices
- Device: where streams of instructions (or kernels) execute (aka “compute device”)
 - can be CPU, GPU, DSP, or any other processor
 - further divided into compute units
 - compute units divided into one or more processing elements (PEs)
 - computations occur within PEs



Execution model

- OpenCL application consists of:
 - **host program**
 - collection of one or more **kernels**
- Host program runs on host
 - OpenCL does not define details of how host program works, only how it interacts with objects defined within OpenCL
- Kernels execute on OpenCL devices
 - Do real work of application
 - Typically simple functions that transform input memory objects into output memory objects
 - **OpenCL kernels**: functions written in OpenCL C
 - **Native kernels**: functions created outside OpenCL (function pointer) [OPTIONAL]

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Execution model

- The OpenCL execution model defines how kernels execute

Assignment Project Exam Help

- How do individual kernels run on a device?

<https://powcoder.com>

- How does the host define the context for kernel execution?

Add WeChat powcoder

- How are the kernels enqueued for execution?

Kernel execution

- Host program issues **command** that submits kernel for execution on device
- Runtime system creates an integer index space
 - Instance of kernel executes for each point in this index space
- Each instance of an executing kernel: **work-item**
 - identified by coordinates in index space
 - coordinates are global ID for work-item
- Command creates collection of work-items, each of which uses same sequence of instructions defined by single kernel
- Sequence of instructions same, but behavior of each work-item can vary (branch statements or data selected through global ID)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Kernel execution

- Work-items organized into **work-groups**
- Provide coarse-grained decomposition of index space
- Exactly span global index space
- Work-groups same size in corresponding dimensions, and this size evenly divides global size in each dimension
- Work-groups assigned unique ID with same dimensionality as index space of work-items
- Work-items assigned unique local ID within work-group: can be uniquely identified by its global ID *or* by a combination of its local ID and work-group ID

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Kernel execution

- Work-items in given work-group execute concurrently on PEs of single compute unit

Assignment Project Exam Help

- Implementation may serialize execution of kernels (may even serialize execution of workgroups in single kernel invocation)

<https://powcoder.com>

Add WeChat powcoder

- OpenCL only assures that workitems within a work-group execute concurrently
- You can never assume that work-groups or kernel invocations execute concurrently

Kernel execution

- Index space spans an N -dimensioned range of values (**NDRange**)

Assignment Project Exam Help

- N can be 1, 2, or 3

<https://powcoder.com>

- Integer array of length N specifying size of index space in each dimension

Add WeChat powcoder

- Work-item's global and local ID is an N -dimensional tuple
- Work-groups assigned IDs using a similar approach to that used for work-items

Kernel execution

$$L_x = G_x / W_x$$

$$L_y = G_y / W_y$$

$$g_x = w_x * L_x + l_x$$

$$g_y = w_y * L_y + l_y$$

$$w_x = g_x / L_x$$

$$w_y = g_y / L_y$$

$$l_x = g_x \% L_x$$

$$l_y = g_y \% L_y$$

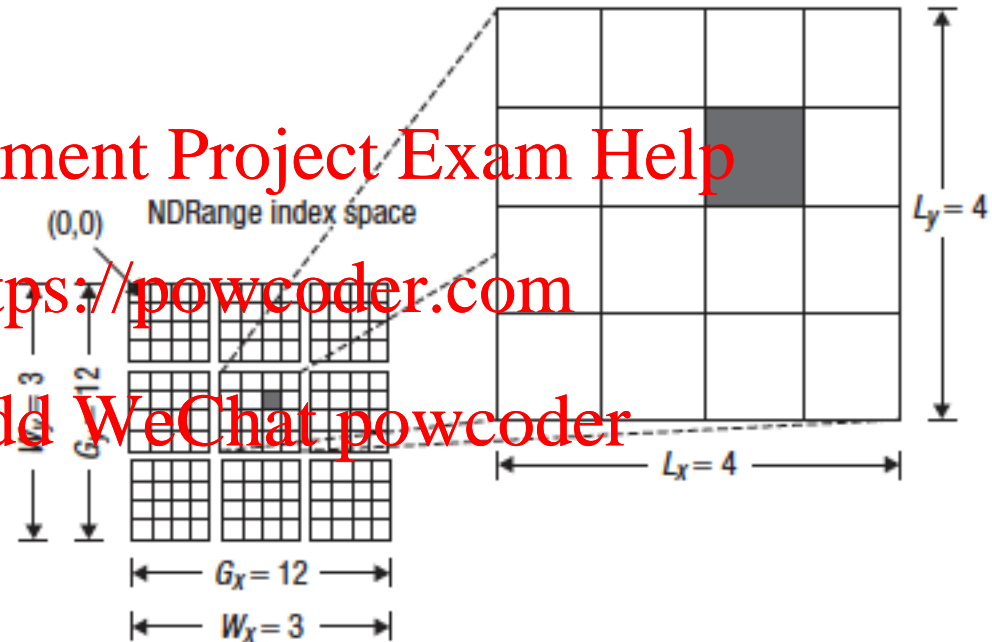
$$g_x = w_x * L_x + l_x + o_x$$

$$g_y = w_y * L_y + l_y + o_y$$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Context

- In OpenCL, computation takes place on device
- But host:
 - Defines and establishes context kernels
 - Defines NDRanges
 - Defines queues that control details of how/when kernels execute
- Context defines environment within which kernels are defined and execute:
 - **Devices:** collection of OpenCL devices to be used by host
 - **Kernels:** OpenCL functions that run on devices
 - **Program objects:** program source code and executables that implement kernels
 - **Memory objects:** set of objects in memory that are visible to OpenCL devices and contain values that can be operated on by instances of a kernel


Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Context

- Created and manipulated by host using the OpenCL API
- Context also contains one or more “program objects”
 - think of these as a dynamic library from which the functions used by the kernels are pulled
- Host program defines devices within context: only at that point is it possible to know how to compile the program source code to create the code for the kernels



Context

- Program object built at runtime within host program (like shader program)
- Context also defines how kernels interact with memory
- On heterogeneous platform, often multiple address spaces to manage
- Devices may have range of different memory architectures
- OpenCL introduces idea of “memory objects”
 - explicitly defined on host
 - explicitly moved between host and OpenCL devices
 - extra burden on programmer, but allows support for much wider range of platforms

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Command-Queues

- Interaction between host and devices occurs through commands posted by host to **command-queue**
- Created by host and attached to single device after context has been defined
- Host places commands into command-queue, and commands are then scheduled for execution on the associated device
- OpenCL supports three types of commands:
 - **Kernel execution commands** execute kernel on PEs of device
 - **Memory commands** transfer data between host and different memory objects, move data between memory objects, or map and unmap memory objects from host address space
 - **Synchronization commands** put constraints on order of execution

Command-Queues

Typical host program

- Define context and command-queues
- Define memory and program objects
- Builds data structures needed on host
- Use command-queue to move memory objects from the host to devices
- Attach kernel arguments to memory objects
- Submit kernels to command-queue for execution
- When kernel completed, memory objects copied back to host

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Command-Queues

- When multiple kernels are submitted, may need to interact
 - E.g., one set of kernels may generate memory objects that a following set of kernels needs to manipulate
 - Synchronization commands can be used to force first set to complete before following set begins
- Many additional subtleties associated with how the commands work in OpenCL
- Commands always execute asynchronously to host
- Host submits commands to command-queue and continues without waiting for them to finish
- If necessary for host to wait on command, can be explicitly established with synchronization command

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Command-Queues

Commands within single queue execute relative to each other in one of 2 modes:

- **In-order execution:** Commands launched in order in which they appear in command-queue and complete in order (serializes execution order of commands)
- **Out-of-order execution:** Commands issued in order but do not wait to complete before the following commands execute (order constraints enforced by programmer through explicit synchronization mechanisms) [OPTIONAL]

Command-Queues

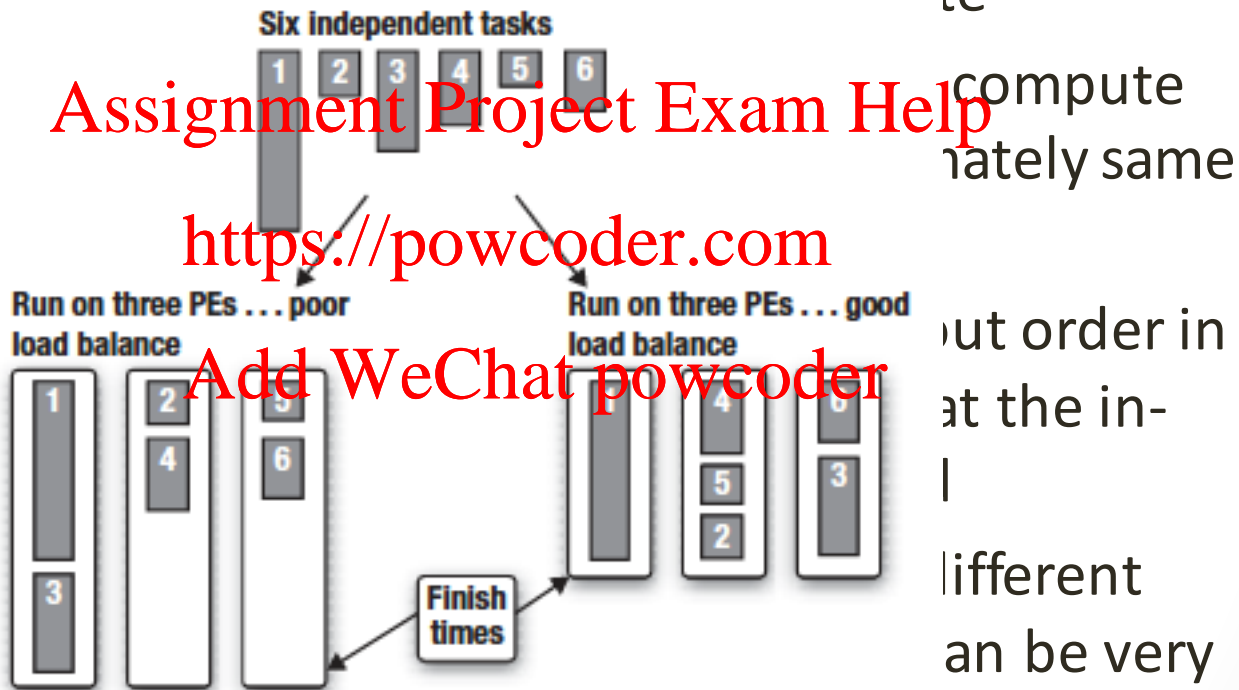
- Why out-of-order? Remember load balancing?
- Application is done until all of kernels complete

- Efficient | units to k amount c

- You can c which yo order exe

- But what amounts

hard! Out-of-order queue can take care of this for you



Command-Queues

- **Automatic load balancing:** Commands can execute in any order, so if compute unit finishes its work early, it can immediately fetch a new command and start executing new kernel
- Commands generate event objects
 - Command can be told to wait until certain conditions on event objects exist
 - Events can also be used to coordinate execution between host and devices
- Also possible to associate multiple queues with single context for any devices within that context
 - Run concurrently and independently with no explicit mechanisms within OpenCL to synchronize between them

Memory Model

Two types of memory objects

- **Buffer object:**

- contiguous block of memory made available to kernels
- programmer can map data structures onto this buffer and access buffer through pointers
- flexibility to define just about any data structure

- **Image object:**

- restricted to holding images
- storage format may be optimized to needs of specific device
- important to give an implementation freedom to customize image format
- opaque object
- OpenCL provides functions to manipulate images, but other than these specific functions, the contents of image object are hidden from kernel program

Assignment Project Exam Help

<https://powcoder.com>

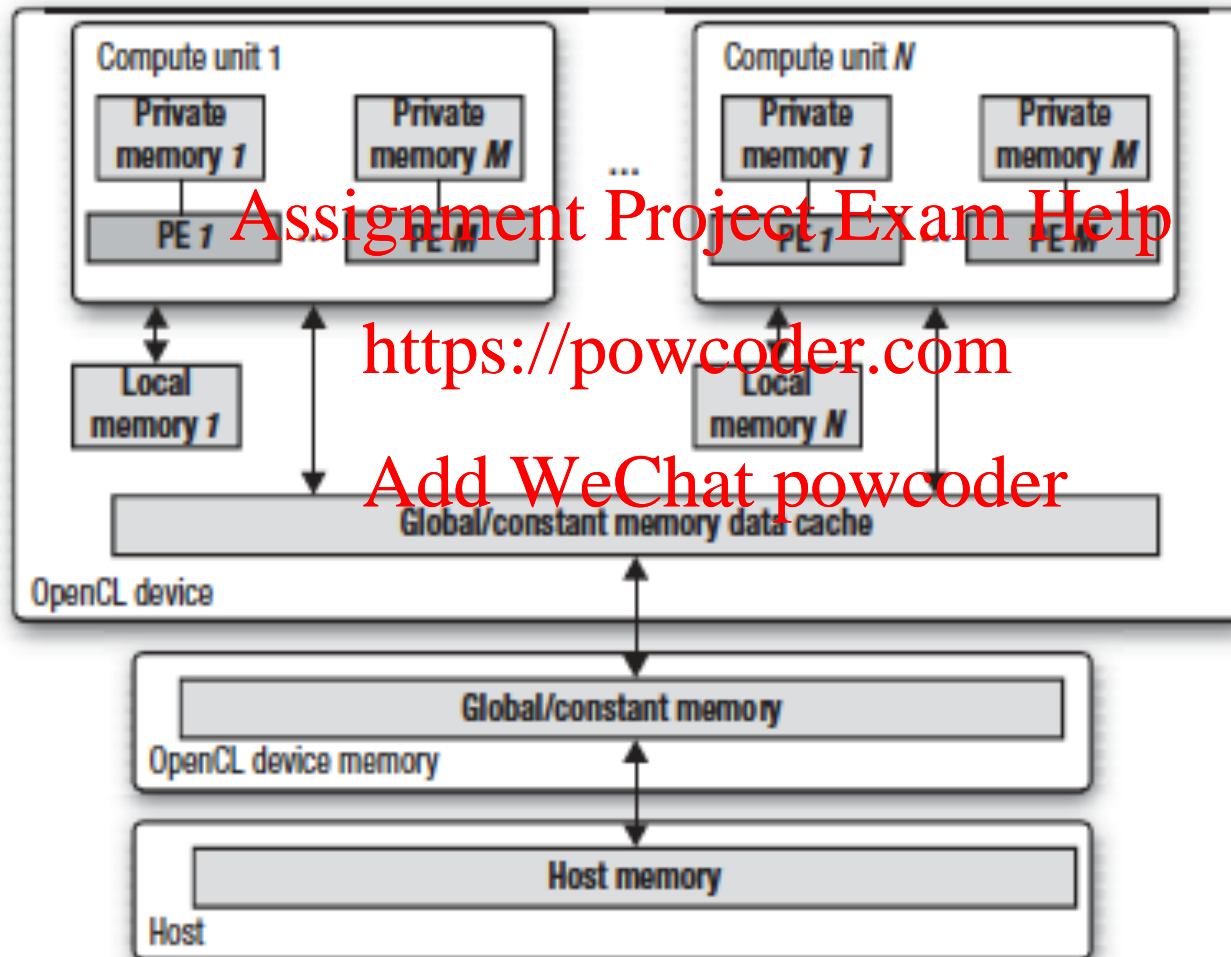
Add WeChat powcoder

Memory Model

OpenCL memory model defines five distinct memory regions:

- **Host memory:** visible only to host
- **Global memory:** permits read/write access to all work-items in all work-groups
- **Constant memory:** region of global memory that remains constant during kernel execution
 - host allocates and initializes
 - work-items have read-only access
- **Local memory:** local to work-group
 - can be used to allocate variables shared by all work-items
 - may be implemented as dedicated regions of memory on device or mapped onto sections of global memory
- **Private memory:** private to work-item
 - not visible to other work-items

Memory Model



Programming Model

- How to map parallel algorithms onto OpenCL
- Programming models intimately connected to how programmers reason about their algorithms
- OpenCL defined with two different programming models in mind: task parallelism and data parallelism
- Also possible to think in terms of a hybrid model: tasks that contain data parallelism

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Programming Model

Data-Parallel Programming Model

- Problems well suited are organized around data structures, the elements of which can be updated concurrently
- Single logical sequence of instructions applied concurrently to elements of data structure
- Structure of algorithm is designed as sequence of concurrent updates to data structures within problem
- Natural fit with OpenCL's execution model
- Key is the NDRange defined when kernel is launched
- Algorithm designer aligns data structures with NDRange index space and maps them onto OpenCL memory objects
- Kernel defines sequence of instructions to be applied concurrently as work-items

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Programming Model

Data-Parallel Programming Model

- Work-items in single work-group may need to share data (local memory region)

Assignment Project Exam Help

<https://powcoder.com>

- Regardless of order in which work-items complete, same results should be produced

Add WeChat powcoder

- Work-items in same work-group can participate in a **work-group barrier** (all must execute before any continuing)
- NB: no mechanism for synchronization between work-items from different work-groups

Programming Model

Data-Parallel Programming Model

- **Single Instruction Multiple Data** or SIMD: no branch statements in kernel, each work-item will execute identical operations but on subset of data items selected by its global ID
- **Single Program Multiple Data** or SPMD: branch statements within a kernel leading each work-item to possibly execute very different operations
- On platforms with restricted bandwidth to instruction memory or if PEs map onto vector unit, SIMD model can be dramatically more efficient

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Programming Model

Data-Parallel Programming Model

- Vector instructions strictly SIMD
- E.g., numerical integration program ($4.0/(1+x^2)$)

Assignment Project Exam Help

```
float8 x, psum_vec;  
float8 ramp = ((float8) {0.5, 1.5, 2.5, 3.5,  
4.5, 5.5, 6.5, 7.5});  
float8 four = ((float8) {4.0, 4.0, 4.0, 4.0, 4.0, 4.0, 4.0, 4.0});  
float8 one = ((float8) {1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0, 1.0});  
float step_number; // step number from loop index  
float step_size; // Input integration step size  
. . . and later inside a loop body . . .  
x = ((float8)step_number + ramp)*step_size;  
psum_vec+=four/(one + x*x);
```

Programming Model

Task-Parallel Programming Model

- Task = kernel that executes as a single work-item regardless of NDRange used by other kernels in application
- Concurrency is internal to the task (eg, vector operations on vector types)
- Kernels submitted as tasks that execute at the same time with an out-of-order queue
- Tasks connected into task graph using OpenCL's event model
 - Commands submitted to event queue may optionally generate events
 - Subsequent commands can wait for these events before executing

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder